

# Neural Networks & Deep Learning: ICP3

**Name: Pavani Medavarthi**

**Student ID: 700741643**

**GitHub Link:** [https://github.com/Pavanimedavarthi/NN-DL\\_Summer-2](https://github.com/Pavanimedavarthi/NN-DL_Summer-2)

**Video Link:** [https://drive.google.com/file/d/14nBpHMEv0U1HJpq-nSeQJ3XTyE-T1d\\_R/view?usp=drive\\_link](https://drive.google.com/file/d/14nBpHMEv0U1HJpq-nSeQJ3XTyE-T1d_R/view?usp=drive_link)

**1. Follow the instruction below and then report how the performance changed.(apply all at once)**

- Convolutional input layer, 32 feature maps with a size of  $3 \times 3$  and a rectifier activation function.
- Dropout layer at 20%.
- Convolutional layer, 32 feature maps with a size of  $3 \times 3$  and a rectifier activation function.
- Max Pool layer with size  $2 \times 2$ .
- Convolutional layer, 64 feature maps with a size of  $3 \times 3$  and a rectifier activation function.
- Dropout layer at 20%.
- Convolutional layer, 64 feature maps with a size of  $3 \times 3$  and a rectifier activation function.
- Max Pool layer with size  $2 \times 2$ .
- Convolutional layer, 128 feature maps with a size of  $3 \times 3$  and a rectifier activation function.
- Dropout layer at 20%.
- Convolutional layer, 128 feature maps with a size of  $3 \times 3$  and a rectifier activation function.

- Max Pool layer with size  $2 \times 2$ .
- Flatten layer.
- Dropout layer at 20%.
- Fully connected layer with 1024 units and a rectifier activation function.
- Dropout layer at 20%.
- Fully connected layer with 512 units and a rectifier activation function.
- Dropout layer at 20%.
- Fully connected output layer with 10 units and a Softmax activation function

```
import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers.convolutional import Conv2D, MaxPooling2D
from keras.constraints import maxnorm
from keras.utils import np_utils
from keras.optimizers import SGD

# Fix random seed for reproducibility
np.random.seed(7)

# Load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# Normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# One hot encode outputs
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_test.shape[1]
```

```

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu'))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 5
learning_rate = 0.01
decay_rate = learning_rate / epochs
sgd = SGD(lr=learning_rate, momentum=0.9, decay=decay_rate, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))

```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 [=====] - 4s 0us/step  
Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
dropout (Dropout)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
dropout_1 (Dropout)	(None, 16, 16, 64)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
dropout_2 (Dropout)	(None, 8, 8, 128)	0
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dropout_3 (Dropout)	(None, 2048)	0
dense (Dense)	(None, 1024)	2098176
dropout_4 (Dropout)	(None, 1024)	0
dense_1 (Dense)	(None, 512)	524800
dropout_5 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130

=====

Total params: 2,915,114  
Trainable params: 2,915,114  
Non-trainable params: 0

```
/usr/local/lib/python3.10/dist-packages/keras/optimizers/legacy/gradient_descent.py:114: UserWarning: The `lr` argument is deprecated, use
super().__init__(name, **kwargs)
None
Epoch 1/5
1563/1563 [=====] - 542s 346ms/step - loss: 1.8841 - accuracy: 0.3015 - val_loss: 1.6270 - val_accuracy: 0.4133
Epoch 2/5
1563/1563 [=====] - 508s 325ms/step - loss: 1.5111 - accuracy: 0.4500 - val_loss: 1.4114 - val_accuracy: 0.4941
Epoch 3/5
1563/1563 [=====] - 524s 335ms/step - loss: 1.3803 - accuracy: 0.4994 - val_loss: 1.3012 - val_accuracy: 0.5231
Epoch 4/5
1563/1563 [=====] - 500s 320ms/step - loss: 1.3048 - accuracy: 0.5303 - val_loss: 1.2518 - val_accuracy: 0.5453
Epoch 5/5
1563/1563 [=====] - 501s 321ms/step - loss: 1.2419 - accuracy: 0.5521 - val_loss: 1.2099 - val_accuracy: 0.5608
Accuracy: 56.08%
```

Did the performance change?

The model's performance is expected to improve by incorporating additional layers and a higher number of feature maps, resulting in enhanced accuracy. Nonetheless, this improvement is accompanied by the downside of increased model complexity and longer training durations, as mentioned in the instructions.

**2. Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4 images to check whether or not the model has predicted correctly.**

```
# Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])

# Convert the predictions to class labels
predicted_labels = np.argmax(predictions, axis=1)

# Convert the actual labels to class labels
actual_labels = np.argmax(y_test[:4], axis=1)

# Print the predicted and actual labels for the first 4 images
print("Predicted labels:", predicted_labels)
print("Actual labels:", actual_labels)
```

```
1/1 [=====] - 0s 457ms/step
Predicted labels: [3 8 8 0]
Actual labels: [3 8 8 0]
```

**3. Visualize Loss and Accuracy using the history object**

```
import matplotlib.pyplot as plt

# Plot the training and validation loss
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()

# Plot the training and validation accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```

